

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>



DeepLearning.AI

# Model Serving

---

# Welcome



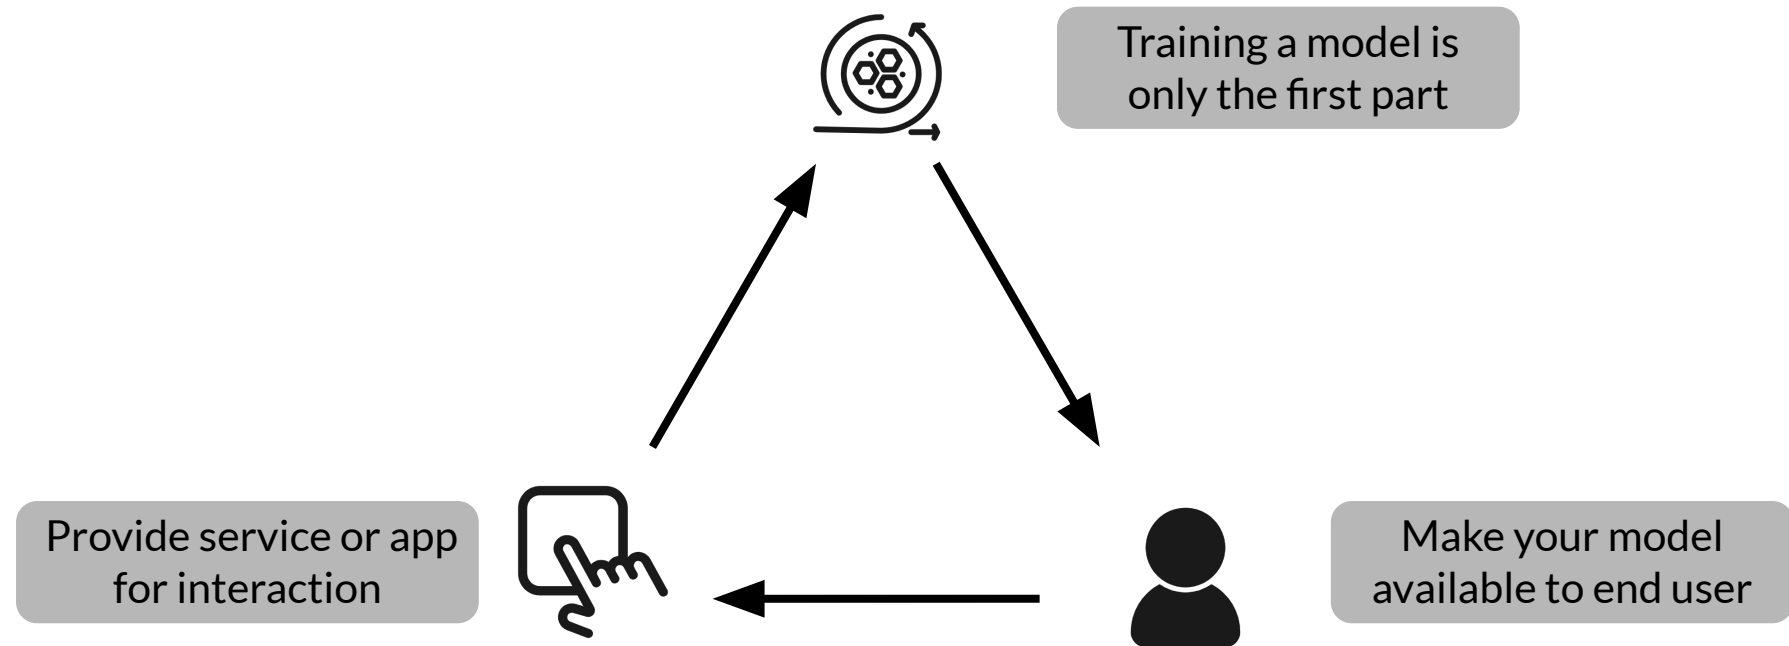
DeepLearning.AI

# Introduction to Model Serving

---

## Introduction

# What exactly is Serving a Model?



# Model Serving Patterns

- A model,
- An interpreter, and
- Input data



Inference

# ML workflows

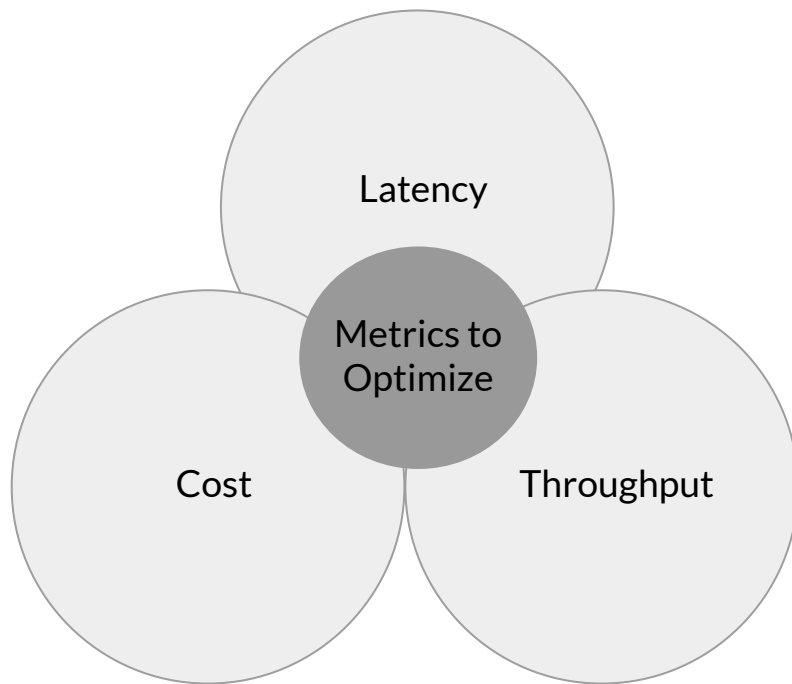
- Model training
- Model prediction



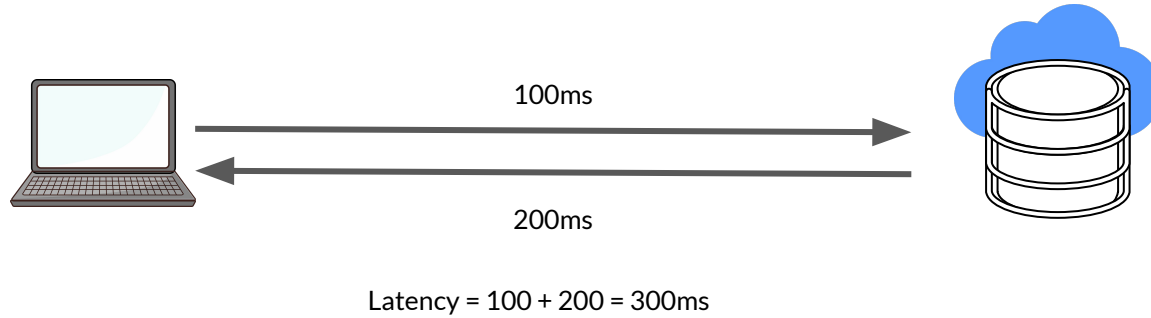
Batch inference

Realtime inference

# Important Metrics



# Latency



- Delay between user's action and response of application to user's action.
- Latency of the whole process, starting from sending data to server, performing inference using model and returning response.
- Minimal latency is a key requirement to maintain customer satisfaction.



# Throughput

- Throughput -> Number of successful requests served per unit time say one second.
- In some applications only throughput is important and not latency.

# Cost

- The cost associated with each inference should be minimised.
  - Important Infrastructure requirements that are expensive:
    - CPU
    - Hardware Accelerators like GPU
    - Caching infrastructure for faster data retrieval.



# Minimizing Latency, Maximizing Throughput

## Minimizing Latency

- Airline Recommendation Service
- Reduce latency for user satisfaction

## Maximizing Throughput

- Airline recommendation service faces high load of inference requests per second.

Scale infrastructure (number of servers, caching requirements etc.) to meet requirements.

# Balance Cost, Latency and Throughput

- Cost increases as infrastructure is scaled
- In applications where latency and throughput can suffer slightly:
  - Reduce costs by GPU sharing
  - Multi-model serving etc.,
  - Optimizing models used for inference





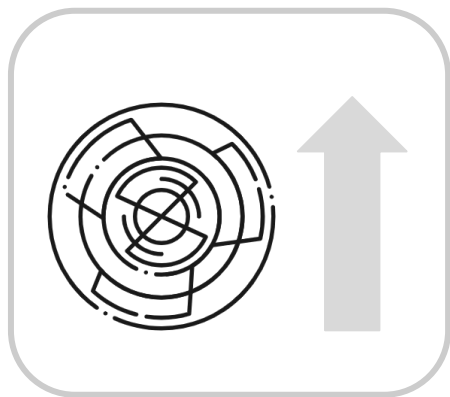
DeepLearning.AI

# Introduction

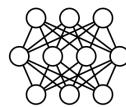
---

## Resources and Requirements for Serving Models

# Optimizing Models for Serving



Model Complexity



Model Size  
Complex functions

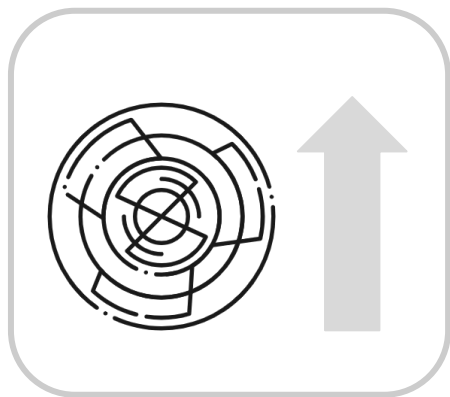


Prediction Latency



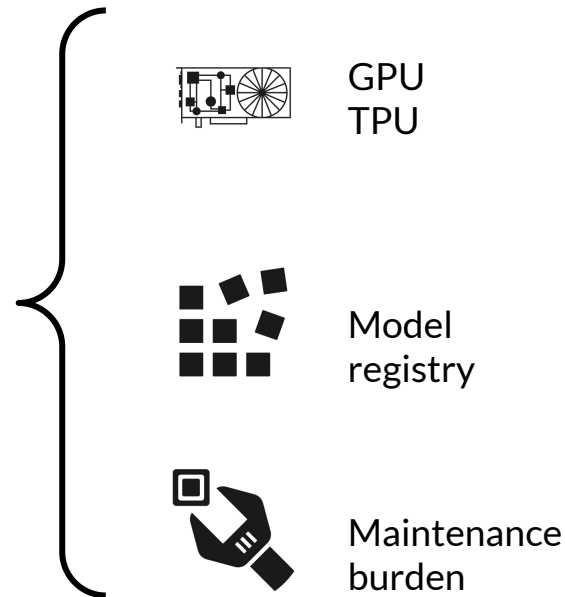
Prediction Accuracy

# As Model Complexity Increases Cost Increases



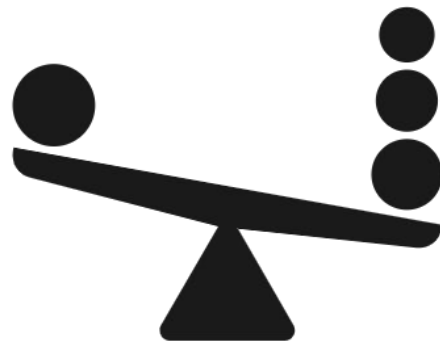
Model Complexity

=



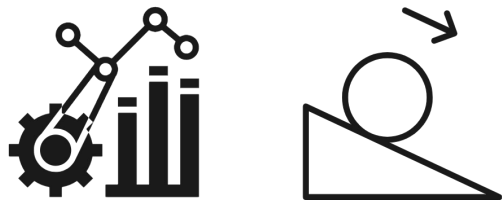
# Balancing Cost and Complexity

The challenge for ML practitioners is to balance complexity and cost.



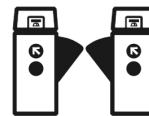


# Optimizing and Satisficing Metrics



Model's optimizing metric:

- Accuracy
- Precision
- Recall



Satisficing (Gating) metric:

- Latency
- Model Size
- GPU load

# Optimizing and Satisficing Metrics

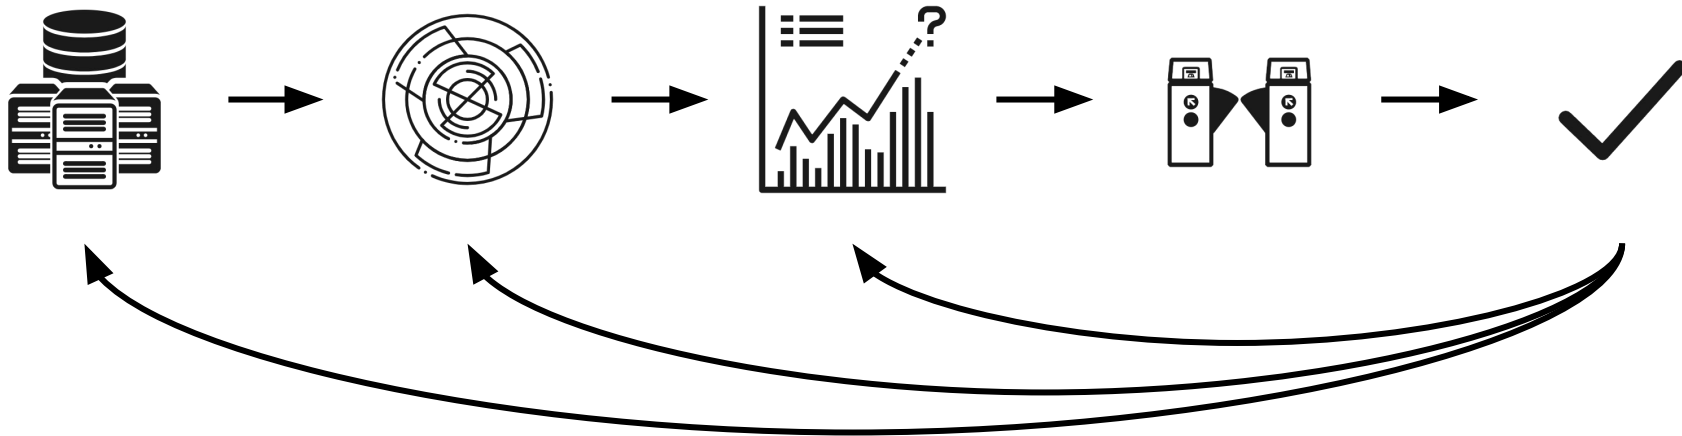
Specify serving infrastructure

Increase model complexity

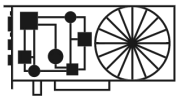
Improve predictive power

Hit gating metrics

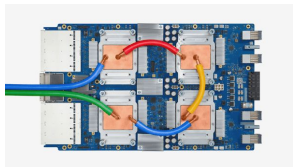
Accept



# Use of Accelerators in Serving Infrastructure



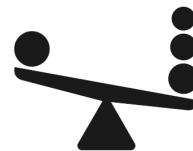
GPUs for parallel  
throughput



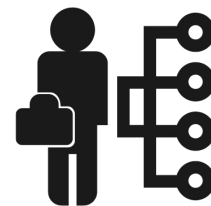
TPUs for complex  
models and large  
batches



Hardware  
choices impact  
cost



Balancing  
complexity and  
hardware choices

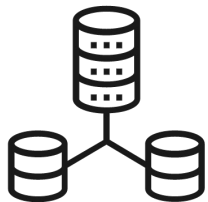


Choices made at  
organizational  
level

# Maintaining Input Feature Lookup

- Prediction request to your ML model might not provide all features required for prediction
- For example, estimating how long food delivery will require accessing features from a data store:
  - Incoming orders (not included in request)
  - Outstanding orders per minute in the past hour
- Additional pre-computed or aggregated features might be read in real-time from a data store
- Providing that data store is a cost

# NoSQL Databases: Caching and Feature Lookup



NoSQL  
Databases

## Google Cloud Memorystore

In memory cache, sub-millisecond read latency

## Google Cloud Firestore

Scaleable, can handle slowly changing data, millisecond read latency

## Google Cloud Bigtable

Scaleable, handles dynamically changing data, millisecond read latency

## Amazon DynamoDB

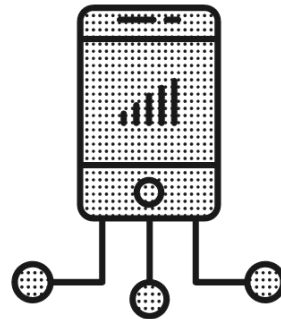
Single digit millisecond read latency, in memory cache available

**Expensive.**  
Carefully choose  
caching  
requirements

# Model Deployments

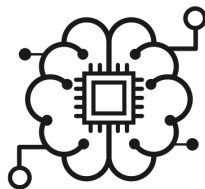
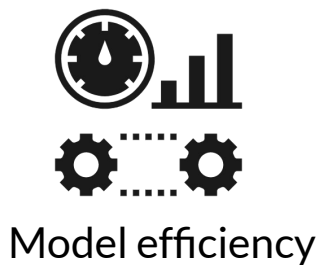


- Huge data centers

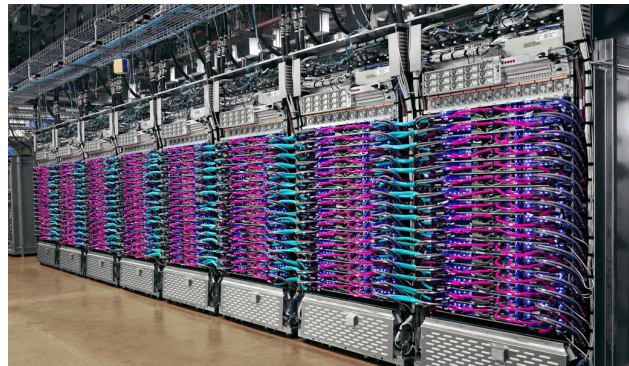


- Embedded devices

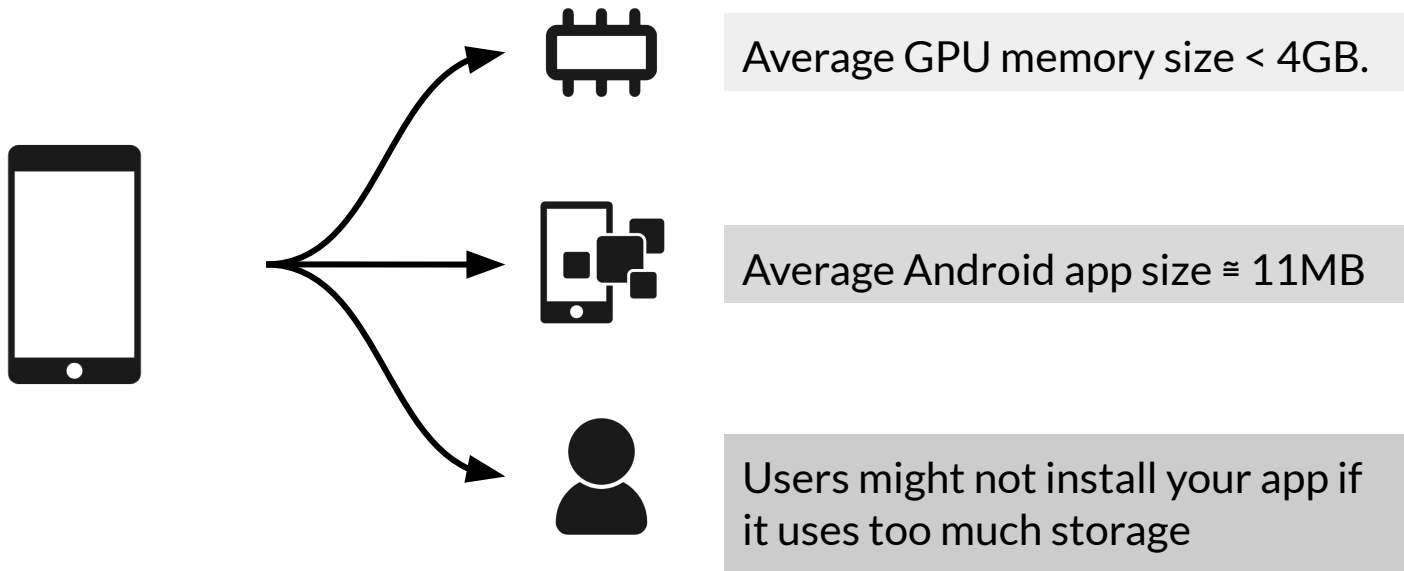
# Running in Huge Data Centers



- Optimize resource utilization
- Reduce cost

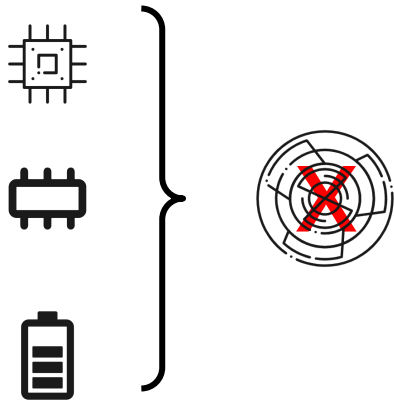


# Constrained Environment: Mobile Phone

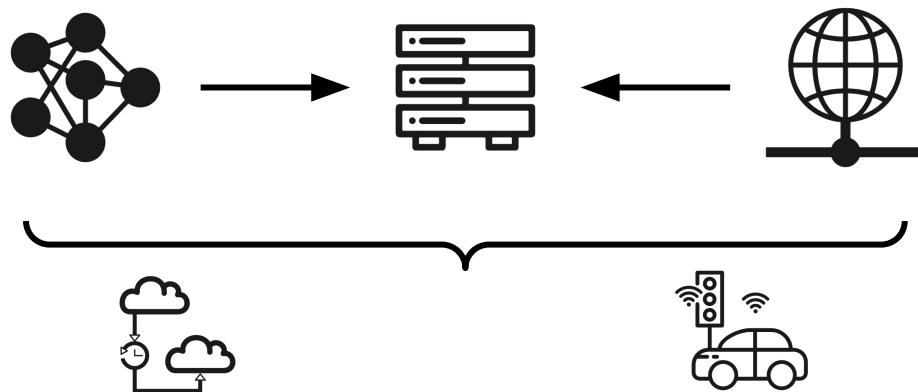




# Restrictions in a Constrained Environment



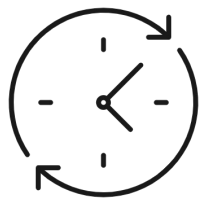
Large, complex models  
cannot be deployed to  
edge devices



Will not work when prediction latency is important. E.g. autonomous car.

# Prediction Latency is Almost Always Important

- Opt for on-device inference whenever possible
  - Enhances user experience by reducing the response time of your app



Millisecond  
turnaround



Model efficiency

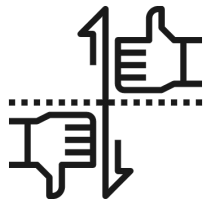


Cost

# Choose Best Model for the Task



# Other Strategies



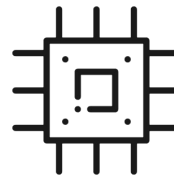
Profile and  
Benchmark



Optimize  
Operators

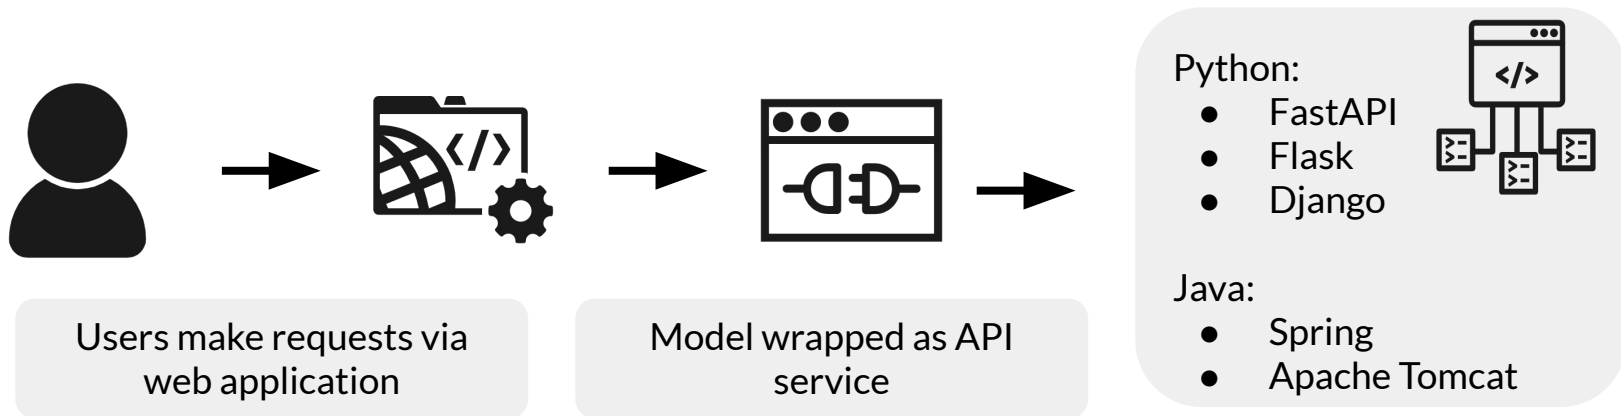


Optimize  
Model



Tweak  
Threads

# Web Applications for Users



# Serving systems for easy deployment



- Centralized model deployment
- Predictions as service



Eliminates need for custom web applications



Deployment just a few lines of code away



Easy to rollback/update models on the fly

# Clipper



Open-source  
project from  
UC Berkeley



Multiple  
modeling  
frameworks



RESTful API



Cluster and  
resources  
management



Settings for  
reliable latency

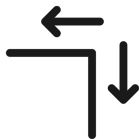
# TensorFlow Serving



Open-source  
project from  
Google



Serve  
TensorFlow  
models easily



Extensible to  
serve other  
model types



Uses REST and  
gRPC protocol



Version  
manager



# Advantages of Serving with a Managed Service



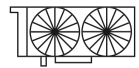
Realtime endpoint for low-latency predictions on massive batches



Deployment of models trained on premises or on the Google Cloud Platform



Scale automatically based on traffic



Use GPU/TPU for faster predictions



DeepLearning.AI

# TensorFlow Serving

---

## Installing and Running TensorFlow Serving

# Install TensorFlow Serving

- Docker Images:
  - Easiest and most recommended method
  - Easiest way to get GPU support with TF Serving

```
docker pull tensorflow/serving
```

```
docker pull tensorflow/serving:latest-gpu
```

# Install TensorFlow Serving

Available Binaries	
<code>tensorflow-model-server</code>	<code>tensorflow-model-server-universal:</code>
<ol style="list-style-type: none"><li>1. Fully optimized server</li><li>2. Uses some platform specific compiler optimizations</li><li>3. May not work on older machines</li></ol>	<ol style="list-style-type: none"><li>1. Compiled with basic optimizations</li><li>2. Doesn't include platform specific instruction sets</li><li>3. Works on most of the machines</li></ol>

# Install TensorFlow Serving

- Building From Source
  - See the complete documentation  
[https://www.tensorflow.org/tfx/serving/setup#building\\_from\\_source](https://www.tensorflow.org/tfx/serving/setup#building_from_source)
- Install using Aptitude (apt-get) on a Debian-based Linux system

# Install TensorFlow Serving

```
!echo "deb http://storage.googleapis.com/tensorflow-serving-apt stable
tensorflow-model-server tensorflow-model-server-universal" | tee
/etc/apt/sources.list.d/tensorflow-serving.list && \
curl
https://storage.googleapis.com/tensorflow-serving-apt/tensorflow-serving.
release.pub.gpg | apt-key add -
!apt update

!apt-get install tensorflow-model-server
```

# Import the MNIST Dataset

```
mnist = tf.keras.datasets.mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
# Scale the values of the arrays below to be between 0.0 and 1.0.
train_images = train_images / 255.0
test_images = test_images / 255.0
```

# Import the MNIST Dataset

```
# Reshape the arrays below.
```

```
train_images = train_images.reshape(train_images.shape[0], 28, 28, 1)
```

```
test_images = test_images.reshape(test_images.shape[0], 28, 28, 1)
```

```
print('\ntrain_images.shape: {}, of {}'.format(train_images.shape,  
train_images.dtype))
```

```
print('test_images.shape: {}, of {}'.format(test_images.shape, test_images.dtype))
```

```
train_images.shape: (60000, 28, 28, 1), of float64
```

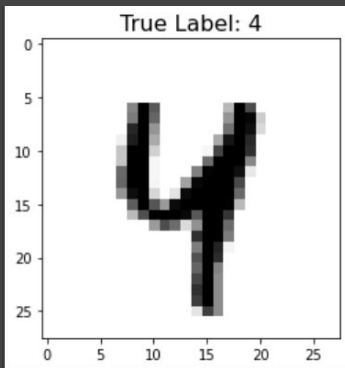
```
test_images.shape: (10000, 28, 28, 1), of float64
```



# Look at a Sample Image

```
idx = 42
```

```
plt.imshow(test_images[idx].reshape(28,28), cmap=plt.cm.binary)  
plt.title('True Label: {}'.format(test_labels[idx]), fontdict={'size': 16})  
plt.show()
```



# Build a Model

```
# Create a model.
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Conv2D(input_shape=(28,28,1), filters=8, kernel_size=3,  
                           strides=2, activation='relu', name='Conv1'),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(10, activation=tf.nn.softmax, name='Softmax')  
)  
  
model.summary()
```

# Train the Model

```
# Configure the model for training.  
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])  
  
epochs = 5  
  
# Train the model.  
history = model.fit(train_images, train_labels, epochs=epochs)
```

# Evaluate the Model

```
# Evaluate the model on the test images.  
results_eval = model.evaluate(test_images, test_labels, verbose=0)  
  
for metric, value in zip(model.metrics_names, results_eval):  
    print(metric + ': {:.3}'.format(value))  
  
loss: 0.098  
accuracy: 0.969
```

# Save the Model

```
MODEL_DIR = tempfile.gettempdir()
version = 1
export_path = os.path.join(MODEL_DIR, str(version))

if os.path.isdir(export_path):
    print('\n Already saved a model, cleaning up\n')
    !rm -r {export_path}

model.save(export_path, save_format="tf")

print('\nexport_path = {}'.format(export_path))
!ls -l {export_path}
```

# Launch Your Saved Model

```
os.environ["MODEL_DIR"] = MODEL_DIR
```

```
%%bash --bg
```

```
nohup tensorflow_model_server \  
  --rest_api_port=8501 \  
  --model_name=digits_model \  
  --model_base_path="${MODEL_DIR}" >server.log 2>&1  
!tail server.log
```

# Send an Inference Request

```
data = json.dumps({"signature_name": "serving_default", "instances":  
test_images[0:3].tolist()})  
  
headers = {"content-type": "application/json"}  
  
json_response =  
    requests.post('http://localhost:8501/v1/models/digits_model:predict',  
                  data=data, headers=headers)  
  
predictions = json.loads(json_response.text)['predictions']
```

# Plot Predictions

```
plt.figure(figsize=(10,15))

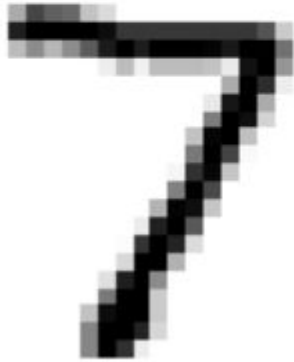
for i in range(3):
    plt.subplot(1,3,i+1)
    plt.imshow(test_images[i].reshape(28,28), cmap = plt.cm.binary)
    plt.axis('off')
    color = 'green' if np.argmax(predictions[i]) == test_labels[i] else 'red'
    plt.title('Prediction: {}\n True Label: {}'.format(np.argmax(predictions[i]),
test_labels[i]), color=color)

plt.show()
```



# Results Demo

Prediction: 7  
True Label: 7



Prediction: 2  
True Label: 2



Prediction: 1  
True Label: 1

